



Log in

Create Account



Sejal Jaiswal
May 20th, 2019

MUST READ

PANDAS

+2

Data Preparation with pandas

In this tutorial, you will learn why it is important to pre-process data and how to do it in pandas.

This tutorial will walk you through some basic concepts and steps for data preparation. Data preparation is the first step after you get your hands on any kind of dataset. This is the step when you pre-process raw data into a form that can be easily and accurately analyzed. Proper data preparation allows for efficient analysis - it can eliminate errors and inaccuracies that could have occurred during the data gathering process and can thus help in removing some bias resulting from poor data quality. Therefore a lot of an analyst's time is spent on this vital step.

Loading data, cleaning data (removing unnecessary data or erroneous data), transforming data formats, and rearranging data are the various steps involved in the data preparation step. In this tutorial, you will work with Python's Pandas library for data preparation.

Make sure you understand concepts like Pandas DataFrame, Series, etc. since you will be reading about it a lot in this tutorial. You can learn more in-depth about pandas in [DataCamp's Pandas tutorial](#).

In this tutorial:

- First, you'll start with a short introduction to [Pandas](#) - the library that is used.
- Then you will load the data.

[Want to leave a comment?](#)

- You will then learn some [data transformation](#) tricks: replacing values, concatenating pandas series, adding knowledge to your dataset using map function, discretizing continuous data, and finally about [dummy variables and one-hot encoding](#).

On the way, there are some tips, tricks, and also reference to other in-depth material in case you want to dive deeper into a topic.

Let's start with a short introduction to **Pandas**...

Pandas

Pandas is a software library written for Python. It is very famous in the data science community because it offers powerful, expressive, and flexible data structures that make data manipulation, analysis easy AND it is freely available. Furthermore, if you have a specific and new use case, you can even share it on one of the Python mailing lists or on [pandas GitHub site](#) - in fact, this is how most of the functionalities in pandas have been driven, by real-world use cases.

For an excellent introduction to pandas, be sure to check out [DataCamp's Pandas Foundation course](#). It is very interactive, and the first chapter is free!

To use the pandas library, you need to first import it. Just type this in your python console:

```
import pandas as pd
```

Loading data

The first step for data preparation is to... well, get some data. If you have a .csv file, you can easily load it up in your system using the `.read_csv()` function in pandas. You can then type:

[Want to leave a comment?](#)

Missing Data

Handling Missing Data

This is a widespread problem in the data analysis world.

Note: If your dataset completely lacks data from a specific scenario, or say does not cover a certain population... well, this is not what the topic covers. In fact, this is a more core error, and one *uhmm uhmm* efficient way of correcting it is by re-evaluating the data collection process.

Missing data can

arise in the dataset due to multiple reasons: the data for the specific field was not added by the user/data collection application, data was lost while transferring manually, a programming error, etc. It is sometimes essential to understand the cause because this will influence how you deal with such data. But this is coming up later on in this tutorial. For now, let's focus on what to do if you have missing data...

For numerical data, pandas uses a floating point value **NaN** (Not a Number) to represent missing data. It is a unique value defined under the library **Numpy** so we will need to import it as well. NaN is the default missing value marker for reasons of computational speed and convenience. This is a sentinel value, in the sense that it is a dummy data or flag value that can be easily detected and worked with using functions in pandas.

Tip: If you want to consider infinity(`inf`) and -infinity (`-inf`) to be missing value in computations, you can set `pandas.options.mode.use_inf_as_na = True`.

Let's play with some

data...

```
import numpy as np

# Creating a pandas series
data = pd.Series([0, 1, 2, 3, 4, 5, np.nan, 6, 7, 8])
```

Want to leave a comment?

```
0    False
1    False
2    False
3    False
4    False
5    False
6     True
7    False
8    False
9    False
dtype: bool
```

Above, we used the function `isnull()` which returns a boolean true or false value. True, when the data at that particular index is actually missing or NaN. The opposite of this is the `notnull()` function.

```
# To check where the dataset does not contain null value - opposite of isnull()
data.notnull()

0     True
1     True
2     True
3     True
4     True
5     True
6    False
7     True
8     True
9     True
dtype: bool
```

Furthermore, we can use the `dropna()` function to filter out missing data and to remove the null (missing) value and see only the non-null values. However, the NaN value is not really deleted and can still be found in the original dataset.

[Want to leave a comment?](#)

```
0    0.0
1    1.0
2    2.0
3    3.0
4    4.0
5    5.0
7    6.0
8    7.0
9    8.0
dtype: float64
```

data

```
0    0.0
1    1.0
2    2.0
3    3.0
4    4.0
5    5.0
6    NaN
7    6.0
8    7.0
9    8.0
dtype: float64
```

What you can do to really "drop" or delete the NaN value is either store the new dataset (without NaN) so that the original data Series is not tampered or apply a drop **inplace**. The **inplace** argument has a default value of false.

```
not_null_data = data.dropna()
not_null_data
```

```
0    0.0
1    1.0
```

[Want to leave a comment?](#)

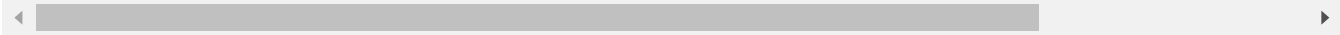
```
8    7.0
9    8.0
dtype: float64
```

```
# Drop the 6th index in the original 'data' since it has a NaN place
data.dropna(inplace = True)
data
```

```
0    0.0
1    1.0
2    2.0
3    3.0
4    4.0
5    5.0
7    6.0
8    7.0
9    8.0
dtype: float64
```

However, dataframes can be more complex and be 2 dimensions, meaning they contain rows and columns. Pandas still has you covered...

```
# Creating a dataframe with 4 rows and 4 columns (4*4 matrix)
data_dim = pd.DataFrame([[1,2,3,np.nan],[4,5,np.nan,np.nan],[7,np.nan,np.nan,np.nan],[np.na
data_dim
```



	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	NaN	NaN
2	7.0	NaN	NaN	NaN

[Want to leave a comment?](#)

the drop is not happening **inplace**, so the real dataset is not tampered. Pay attention to the arguments passed to the `dropna()` function to determine how you drop the missing data.

```
# Drop all rows and columns containing NaN value
data_dim.dropna()
```

	0	1	2	3
--	---	---	---	---

```
# Drop all rows and columns containing entirely of NaN value
data_dim.dropna(how = 'all')
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	NaN	NaN
2	7.0	NaN	NaN	NaN

```
# Drop only columns that contain entirely NaN value
# Default is 0 - which signifies rows
data_dim.dropna(axis = 1, how = 'all')
```

	0	1	2
0	1.0	2.0	3.0
1	4.0	5.0	NaN
2	7.0	NaN	NaN
3	NaN	NaN	NaN

[Want to leave a comment?](#)

	0	1.0	2.0
0		4.0	5.0
1		7.0	NaN
2		NaN	NaN

```
# Drop all rows that have more than 2 NaN values
data_dim.dropna(thresh = 2)
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	NaN	NaN

Now you know how to identify and drop missing values - whether to simply see the resultant dataset or do an **inplace** deletion. In many cases, simply dropping the null values is not a feasible option, and you might want to fill in the missing data with some other value. Let's see how you can do that...

Filling in Missing Data

To replace or rather "fill in" the null data, you can use the `fillna()` function. For example, let's try to use the same dataset as above and try to fill in the NaN values with 0.

```
# Check what the dataset looks like again
data_dim
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	NaN	NaN

[Want to leave a comment?](#)

3	NaN	NaN	NaN	NaN
---	-----	-----	-----	-----

```
# Fill the NaN values with 0
data_dim_fill = data_dim.fillna(0)
data_dim_fill
```

	0	1	2	3
0	1.0	2.0	3.0	0.0
1	4.0	5.0	0.0	0.0
2	7.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0

And like with `dropna()` you can also do many other things depending on the kind of argument you pass. Also a reminder that passing the `inplace = True` argument will make the change to the original dataset.

```
# Pass a dictionary to use different values for each column
data_dim_fill = data_dim.fillna({0: 0, 1: 8, 2: 9, 3: 10})
data_dim_fill
```

	0	1	2	3
0	1.0	2.0	3.0	10.0
1	4.0	5.0	9.0	10.0
2	7.0	8.0	9.0	10.0
3	0.0	8.0	9.0	10.0

You can pass a `method` argument to the `fillna()` function that automatically

[Want to leave a comment?](#)

```
data_dim_fill = data_dim.fillna(method='ffill')
data_dim_fill
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	3.0	NaN
2	7.0	5.0	3.0	NaN
3	7.0	5.0	3.0	NaN

You can also limit the number of fills above. For example, fill up only two places in the columns... Also, if you pass **axis = 1** this will fill out row value accordingly.

```
# Pass method to determine how to fill-up the column - forward here
data_dim_fill = data_dim.fillna(method='ffill', limit = 2)
data_dim_fill
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	3.0	NaN
2	7.0	5.0	3.0	NaN
3	7.0	5.0	NaN	NaN

```
# Pass method to determine how to fill-up the row - forward here
data_dim_fill = data_dim.fillna(axis = 1, method='ffill')
data_dim_fill
```

	0	1	2	3
--	---	---	---	---

[Want to leave a comment?](#)

1	4.0	5.0	5.0	5.0
2	7.0	7.0	7.0	7.0
3	NaN	NaN	NaN	NaN

With some understanding of the data and your use-case, you can use the `fillna()` function in many other ways than simply filling it with numbers. You could fill it up using the mean using the `mean()` or the median value `median()` as well...

```
# Check the data_dim dataset
data_dim
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	NaN	NaN
2	7.0	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN

```
# Fill the NaN value with mean values in the corresponding column
data_dim_fill = data_dim.fillna(data_dim.mean())
data_dim_fill
```

	0	1	2	3
0	1.0	2.0	3.0	NaN
1	4.0	5.0	3.0	NaN
2	7.0	3.5	3.0	NaN

[Want to leave a comment?](#)

Replacing Values

So far, you have only worked with missing data (NaN), but there could be situations where you would want to replace a non-null value with a different value. Or maybe a null value is recorded as a random number, and hence needs to be processed as NaN rather than a number. This is where the `replace()` function comes in handy...

Let's create a different dataset this time.

```
data = pd.Series([1,2,-99,4,5,-99,7,8,-99])
data
```

```
0    1
1    2
2   -99
3    4
4    5
5   -99
6    7
7    8
8   -99
dtype: int64
```

```
# Replace the placeholder -99 as NaN
data.replace(-99, np.nan)
```

```
0    0.0
1    1.0
2    2.0
3    3.0
4    4.0
5    5.0
7    6.0
8    7.0
```

[Want to leave a comment?](#)

series and then concatenate the original data series with the new series and then apply the multiple value replace function.

Concatenating Pandas Series

To do this, we can use the `concat()` function in pandas. To continue the indexing after applying the concatenation, you can pass the `ignore_index = True` argument to it.

```
# Create a new Series
new_data = pd.Series([-100, 11, 12, 13])
combined_series = pd.concat([data, new_data], ignore_index = True)
combined_series
```

```
0      1
1      2
2    -99
3      4
4      5
5    -99
6      7
7      8
8    -99
9   -100
10     11
11     12
12     13
dtype: int64
```

```
# Let's replace -99 and -100 as NaN in the new combined_series
data_replaced = combined_series.replace([-99, -100], np.nan)
data_replaced
```

```
0      1.0
1      2.0
```

[Want to leave a comment?](#)

```
7      8.0
8      NaN
9      NaN
10     11.0
11     12.0
12     13.0
dtype: float64
```

```
# Argument passed can also be a dictionary with separate values
data_replaced = combined_series.replace({-99: np.nan, -100: 0})

# Same as: new_data.replace([-99, -100], [np.nan, 0])
data_replaced
```

```
0      1.0
1      2.0
2      NaN
3      4.0
4      5.0
5      NaN
6      7.0
7      8.0
8      NaN
9      0.0
10     11.0
11     12.0
12     13.0
dtype: float64
```

Adding knowledge - Map Function

In some situations, you might want to add more insight to what you already have based on some logic. This is when you can take help from maps and a combination of functions to achieve what you want. The logic can get more complicated, but try understanding with the help of this example

[Want to leave a comment?](#)

```
digits = [0,1,2,3,4,5]
data_number
```

	english	digits
0	zero	0
1	one	1
2	two	2
3	three	3
4	four	4
5	five	5

Let's say you now want to add another column indicating multiples of two as 'Yes' and the rest as 'No'. You can write down a mapping of each distinct English calls to it's corresponding 'Yes' or 'No'.

```
english_to_multiple = {
    'two': 'yes',
    'four': 'yes'
}
```

Then you can call the map function to add the column for when the English column is a multiple of 2. What gets filled in the other non-multiple number columns? Let's find out...

```
data_number['multiple'] = data_number['english'].map(english_to_multiple)
data_number
```

	english	digits	multiple
0	zero	0	NaN

[Want to leave a comment?](#)

2	two	2	yes
3	three	3	NaN
4	four	4	yes
5	five	5	NaN

The other columns are filled with NaN values, and you already know how to further work with missing values. This was a simple example. But hope you can find inspiration to use ideas from here to use the map function to do more stuff and utilize it in your specific use-cases.

Discretization - Cut Function

Sometimes you might want to categorize based on some logic and put all the data into discrete buckets or bins for analysis purpose. You can use the `cut()` function for this. For example, let's first create a dataset containing 30 random numbers between 1 - 100.

```
import random
```

```
data = random.sample(range(1, 101), 30)
```

```
# data
```

```
[45,  
 39,  
 25,  
 83,  
 27,  
 6,  
 73,  
 43,  
 36,  
 93,  
 97,
```

[Want to leave a comment?](#)


```
31,  
22,  
65,  
30,  
3,  
26,  
91,  
12,  
52,  
76,  
63,  
84,  
59,  
53]
```

Say we want to categorize those in terms of some bucket we define ourselves: numbers between 1 - 25, then 25 - 35, 40 - 60 and then 60 - 80 and then the rest. So we define a bucket...

Then we will use the `cut` function.

```
# Defining the starting value for each bucket
```

```
bucket = [1, 25, 35, 60, 80, 100]
```

```
cut_data = pd.cut(data, bucket)
```

```
cut_data
```

```
[(35, 60], (35, 60], (1, 25], (80, 100], (25, 35], ..., (60, 80], (60, 80], (80, 100], (35,  
Length: 30
```

```
Categories (5, interval[int64]): [(1, 25] < (25, 35] < (35, 60] < (60, 80] < (80, 100]]
```

`cut()` is a very useful function, and there is much more you can do with it. It is highly recommended that you check out the [Pandas document](#) to learn more about it.

[Want to leave a comment?](#)

particularly handy, especially when doing machine learning modeling, where the concept of one-hot encoding is famous. Using more technical words: one-hot encoding is the process of converting categorical values into a 1-dimensional numerical vector.

One way of doing this using pandas is to use the `get_dummies()` function. If a column in your dataframe has 'n' distinct values, the function will derive a matrix with 'n' columns containing all 1s and 0s. Let's see this with an example to grasp the concept better...

```
# Creating a DataFrame consisting individual characters in the list
data = pd.Series(list('abcdabababcdabcd'))
data
```

0	a
1	b
2	c
3	d
4	a
5	b
6	a
7	b
8	c
9	d
10	a
11	b
12	c
13	d

```
dtype: object
```

Let's say now you want to have individual vectors indicating the appearance of each character to feed it to a function.

Something like this: for 'a' = [1,0,0,0,1,0,1,0,0,0,1,0,0,0] where 1 is in the position where 'a' exists and 0 for where it does not. Using the `get_dummies()` function will make the task easier.

[Want to leave a comment?](#)

0	1	0	0	0
1	0	1	0	0
2	0	0	1	0
3	0	0	0	1
4	1	0	0	0
5	0	1	0	0
6	1	0	0	0
7	0	1	0	0
8	0	0	1	0
9	0	0	0	1
10	1	0	0	0
11	0	1	0	0
12	0	0	1	0
13	0	0	0	1

However, this is a simple example to help you get started. In real world applications the membership of a character can belong to various categories at the same time, thus you will have to learn more sophisticated methods. Also for much larger data, this function might not be very efficient in terms of speed. DataCamp's [Handling Categorical Data in Python](#) discusses such cases and the topic more in-depth.

What Now?

Well, you are at the end of the tutorial, and you have learned some basic tools to start the

[Want to leave a comment?](#)

Real world use cases can often be more complex and might need its own unique solution, but the idea is to start small and learn on the way and to explore more and get your hands dirty with different datasets! Be sure to check out [DataCamp's Cleaning Data in Python](#) course to learn more tools and tricks and in the end, you get to work on a real-world, messy dataset obtained from the Gapminder Foundation. Go get started!


26
3

COMMENTS

Ryu Yan

16/06/2019 11:53 AM

As a 100% beginner, I've learned a lot from this tutorial, many thx to you!

▲ 1 ↩ [REPLY](#)

DeVon Goldsmith

21/06/2019 03:35 AM

Omg this is golden data!!!

▲ 1 ↩ [REPLY](#)

chetas poojara

27/06/2019 06:29 PM

deep and detailed information, learned a lot, thanks for writing up this informative article.

[free shared passwords](#)

▲ 1 ↩ [REPLY](#)

[Want to leave a comment?](#)

[About](#) [Terms](#) [Privacy](#)

Want to leave a comment?